

# BML lecture #5: Bayesian deep learning

<http://github.com/rbardenet/bml-course>

Julyan Arbel

Statify team, Inria Grenoble Rhône-Alpes & Univ. Grenoble-Alpes, France


The Inria logo is written in a red, cursive script.The Statify logo features a blue wavy line above the word "Statify" in a black, sans-serif font.

- 1** Introduction
- 2** Feed-forward neural networks
- 3** Inference algorithms
- 4** Priors in weight space
- 5** Priors in function space
- 6** Performance metrics


- 1 Introduction**
- 2 Feed-forward neural networks
- 3 Inference algorithms
- 4 Priors in weight space
- 5 Priors in function space
- 6 Performance metrics

What comes to *your* mind when you hear “Bayesian deep learning”?


## Deep neural networks Achilles heels




## Different flavours of neural networks (Jospin et al., 2020)




Point estimate NN



BNN w/ random weights




BNN w/ last-layer rd weights



BNN w/ random activations

- 1 Introduction
- 2 Feed-forward neural networks**
- 3 Inference algorithms
- 4 Priors in weight space
- 5 Priors in function space
- 6 Performance metrics

## Neural networks notations



- ▶ *pre-nonlinearity*  $\mathbf{g}^{(\ell)} = \mathbf{g}^{(\ell)}(\mathbf{x})$ , *post-nonlinearity*  $\mathbf{h}^{(\ell)} = \mathbf{h}^{(\ell)}(\mathbf{x})$

$$\mathbf{g}^{(\ell)}(\mathbf{x}) = \mathbf{W}^{(\ell)} \mathbf{h}^{(\ell-1)}(\mathbf{x}), \quad \mathbf{h}^{(\ell)}(\mathbf{x}) = \varphi(\mathbf{g}^{(\ell)}(\mathbf{x}))$$

- ▶ *nonlinearity* or *activation function*  $\varphi : \mathbb{R} \rightarrow \mathbb{R}$ .
- ▶ *weight matrix*  $\mathbf{W}^{(\ell)}$  of dimension  $H_\ell \times H_{\ell-1}$  including a bias vector



Optimization problem: minimize the loss function

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} \mathcal{L}(\mathbf{w}).$$

With gradient-based optimization:

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \partial_{\mathbf{w}} \mathcal{L}(\mathbf{w}).$$

$\eta > 0$  is a *step size*, or *learning rate*. Gradients are computed as products of gradients between each layer *from right to left*, a procedure called *backpropagation* (Rumelhart, Hinton, and Williams, 1986).

Gradients are approximated on randomly chosen subsets called *batches*: stochastic gradient descent, SGD (Robbins and Monro, 1951). See survey of optimization methods Sun et al. (2019).

- ▶ **Convolutional neural networks (CNN)** are widely used in computer vision.
- ▶ **Recurrent neural networks (RNN)** are advantageous for sequential data, designed to save the output of a layer by adding it back to the input (Hochreiter and Schmidhuber, 1997).
- ▶ **Residual neural networks (ResNet)** have residual blocks which add the output from the previous layer to the layer ahead, so-called *skip-connections* (He et al., 2016). Allows very deep training.

Expressiveness describes neural networks' ability to approximate functions (Cybenko, 1989; Funahashi, 1989; Hornik, Stinchcombe, and White, 1989; Barron, 1994).

### Universal approximation theorem


Neural networks of one hidden layer and suitable activation function can approximate any continuous function on a compact domain, say  $f : [0, 1]^N \rightarrow \mathbb{R}$ , to any desired accuracy.

**But** the size of such networks may be *exponential in the input dimension*  $N$ , which makes them highly prone to overfitting as well as impractical.


Width-depth trade-offs studied by Chatziafratis, Nagarajan, Panageas, and Wang (2020) and Chatziafratis, Nagarajan, and Panageas (2020).

## Classical regime


underfitting



optimum




overfitting



## Generalization and overfitting II

### Modern regime

It was shown recently that when increasing the model size beyond the number of training examples, the model's test error can start *decreasing again* after reaching the interpolation peak: *double-descent* (Belkin et al., 2019).



## Limitations with point-estimate neural networks

- ▶ Inability to distinguish between *in-domain* and *out-of-domain* samples (Lee et al., 2018; Mitros and Mac Namee, 2019; Hein, Andriushchenko, and Bitterwolf, 2019; Ashukha et al., 2020), and the sensitivity to *domain shifts* (Ovadia et al., 2019), which are explained in details later on;
- ▶ Inability to provide reliable uncertainty estimates for a deep neural network's decision and frequently occurring overconfident predictions (Minderer et al., 2021);
- ▶ Lack of transparency and interpretability of a deep neural network's inference model, which makes it difficult to trust their outcomes;
- ▶ Sensitivity to adversarial attacks that make deep neural networks vulnerable for sabotage (Wilson et al., 2016).

- ▶ Uncertainty quantification through the posterior distribution: BNN are shown to be better calibrated than NN
- ▶ Distinguishing between the epistemic uncertainty  $p(\theta|D)$  and the aleatoric uncertainty  $p(y|x, \theta)$ : desirable in small dataset settings, providing high epistemic uncertainty for prediction, avoiding overfitting
- ▶ Integrating prior knowledge: most regularization methods for NN can be understood as setting a prior
- ▶ Interpreting known ML algorithms as approximate Bayesian methods: including regularization, ensembling, constant (learning rate) SGD, etc.

- 1 Introduction
- 2 Feed-forward neural networks
- 3 Inference algorithms**
- 4 Priors in weight space
- 5 Priors in function space
- 6 Performance metrics



Denote data by  $D = \{D_x, D_y\}$  and parameters (weights) by  $\theta$ .

---

**Algorithm 1** Inference procedure for a BNN.

---

Define  $p(\theta|D) = \frac{p(D_y|D_x, \theta)p(\theta)}{\int_{\theta} p(D_y|D_x, \theta')p(\theta')d\theta'}$ ;

**for**  $i = 0$  **to**  $N$  **do**

    Draw  $\theta_i \sim p(\theta|D)$ ;

$y_i = \Phi_{\theta_i}(x)$ ;

**end for**

**return**  $Y = \{y_i | i \in [0, N)\}$ ,  $\Theta = \{\theta_i | i \in [0, N)\}$ ;

---

- ▶ Markov chain Monte Carlo (MCMC), Hamiltonian Monte Carlo (HMC). No-U-Turn sampler (NUTS) is most often used in probabilistic programming languages (Stan, PyMC3, Pyro, etc): it improves over classic HMC by allowing hyperparameters to be set automatically instead of manually
- ▶ Variational inference (VI): scales better than MCMC algorithms. Idea: find an approximate variational distribution in a variational family that is as close as possible to the exact posterior by minimizing the Kullback–Leibler divergence. Turns sampling into optimization.
- ▶ Stochastic variational inference (SVI): scales better than VI, stochastic gradient descent method applied to VI. Gradient of objective is computed only on mini-batches.

### BUT

Stochasticity in gradient estimation stops backpropagation from functioning

### Tricks for Monte Carlo gradient estimation

A number of tricks (see *Monte Carlo Gradient Estimation in Machine Learning*, Mohamed et al, JMLR, 2020):

- ▶ Log-derivative trick: score function estimators
- ▶ Reparameterisation trick: pathwise derivative estimator
- ▶ Measure-valued gradient estimators

- ▶ Bayes-by-backprop (BBB) and probabilistic backpropagation (PBP): implement the reparameterisation trick
- ▶ Monte Carlo dropout: turning dropout into an approximate Bayesian algorithm (variational inference)
- ▶ Bayes via stochastic gradient descent: includes MCMC algorithms based on the SGD dynamic such as stochastic gradient Langevin dynamic (SGLD) and Variational Inference based on SGD dynamic such as ensembling

- 1 Introduction
- 2 Feed-forward neural networks
- 3 Inference algorithms
- 4 Priors in weight space**
- 5 Priors in function space
- 6 Performance metrics

In deep learning, initializing neural networks with appropriate weights is crucial to obtaining convergence.

Adequate initialization can help avoid *vanishing* and *exploding gradients*.

Lottery ticket hypothesis: **frankle2019lottery** proposed an iterative algorithm for parameter pruning in neural networks while saving the original initialization of the weights after pruning, also known as the *winning ticket* of the initialization “lottery”. Neural networks with such winning tickets could outperform unpruned neural networks.

Typical initialization (Glorot initialization): independently sample each bias  $b_i^{(\ell)}$  and each weight  $W_{ij}^{(\ell)}$  from zero-mean Gaussian distributions (Glorot and Bengio, 2010):

$$b_i^{(\ell)} \sim \mathcal{N}\left(0, \sigma_b^{(\ell)}\right), \quad W_{ij}^{(\ell)} \sim \mathcal{N}\left(0, \frac{\sigma_w^{(\ell)}}{H_{\ell-1}}\right),$$

for all  $i = 1, \dots, H_\ell$  and  $j = 1, \dots, H_{\ell-1}$ , where the normalization of weight variances by  $1/H_{\ell-1}$  is conventional to avoid the variance explosion in wide neural networks.

Poole et al. (2016) and Schoenholz et al. (2017) show that there is a critical line, called **Edge of Chaos**, separating signal propagation into two regions in  $(\sigma_b^{(\ell)}, \sigma_w^{(\ell)})$  initialization plane:

$$w_{ij}^{(\ell)} \sim \mathcal{N}\left(0, \frac{\sigma_w^2}{H_{\ell-1}}\right) \text{ and biases } b_i^{(\ell)} \sim \mathcal{N}(0, \sigma_b^2) \text{ for all } \ell, i \text{ and } j.$$

Let

- ▶  $\mathbf{x}_a$  be an input vector of a data point  $a$ .
- ▶  $g_{i,a}^{(\ell)}$  be a pre-activation (centered random variable) at layer  $\ell$  given a data point  $a$ .
- ▶  $q_{aa}^{(\ell)} = \mathbb{E}\left[\left(g_{i,a}^{(\ell)}\right)^2\right]$  the variance of pre-activation at layer  $\ell$  given input  $a$ .
- ▶  $q_{ab}^{(\ell)} = \mathbb{E}\left[g_{i,a}^{(\ell)} g_{i,b}^{(\ell)}\right]$  the covariance between the pre-activations at layer  $\ell$  given two inputs  $a$  and  $b$ .



Two-way recurrence relations:

$$q_{aa}^{(\ell)} = \sigma_w^2 \int \varphi^2 \left( u_1^{(\ell-1)} \right) \mathcal{D}g_{i,a} + \sigma_b^2,$$

$$q_{ab}^{(\ell)} = \sigma_w^2 \int \varphi(u_1^{(\ell-1)}) \varphi(u_2^{(\ell-1)}) \mathcal{D}g_{i,a} \mathcal{D}g_{i,b} + \sigma_b^2,$$

where  $u_1^{(\ell-1)} = \sqrt{q_{aa}^{(\ell-1)}} g_{i,a}$ ,

$u_2^{(\ell-1)} = \sqrt{q_{bb}^{(\ell-1)}} \left( c_{ab}^{(\ell-1)} g_{i,a} + \sqrt{1 - (c_{ab}^{(\ell-1)})^2} g_{i,b} \right)$  and

$c_{ab}^{(\ell)} = \frac{q_{ab}^{(\ell)}}{\sqrt{q_{aa}^{(\ell)}} \sqrt{q_{bb}^{(\ell)}}}$ . Here,  $\mathcal{D}g_{i,a}$  and  $\mathcal{D}g_{i,b}$  stand for the distributions of standard Gaussian pre-activations  $g_{i,a}$  and  $g_{i,b}$ .


For any  $\sigma_w^2$  and  $\sigma_b^2$ , there are limiting points for variance

$q^* = \lim_{\ell \rightarrow \infty} q_{aa}^{(\ell)}$  and for correlation  $c^* = \lim_{\ell \rightarrow \infty} c_{ab}^{(\ell)}$ . Two regions can be defined depending on the value of  $c^*$ : (i) an *ordered* region if  $c^* = 1$ , as any two inputs  $a$  and  $b$ , even far from each others, tend to be fully

correlated in the deep limit  $\ell \rightarrow \infty$ ; (ii) a *chaos* region if  $c^* < 1$ , as any two inputs  $a$  and  $b$ , even close to each others, tend to decorrelate  $\ell \rightarrow \infty$ .

To study whether the point  $c^* = 1$  is *stable*, we need to check the values of the derivative:  $\chi_1 = \left. \frac{\partial c_{ab}^{(\ell)}}{\partial c_{ab}^{(\ell-1)}} \right|_{c_{ab}^{(\ell)}=1}$ . There are three cases: (i) *order*, when  $\chi_1 < 1$ , i.e., the point  $c^* = 1$  is stable; (ii) *transition*, when  $\chi_1 = 1$ ; (iii) *chaos*, when  $\chi_1 > 1$ , i.e., the point  $c^* = 1$  is unstable. Therefore, there is a separating line when  $c^* = 1$  and  $\chi_1 = 1$ . By assigning initialization hyperparameters  $\sigma_w^2$  and  $\sigma_b^2$  according to the separating line, the information propagates as deep as possible from inputs to outputs, until convergence.

## Edge of Chaos IV



**Figure:** (a) Edge of chaos diagram showing the boundary between ordered and chaotic phases as a function of  $\sigma_w^2$  and  $\sigma_b^2$ . (b) The residual  $|q^* - q_{aa}^l|$  as a function of depth on a log-scale with  $\sigma_b^2 = 0.05$  and  $\sigma_w^2$  from 0.01 (red) to 1.7 (purple). Clear exponential behavior is observed. (c) The residual  $|c^* - c_{ab}^l|$  as a function of depth on a log-scale. Again, the exponential behavior is clear. The same color scheme is used here as in (b). From Schoenholz et al. (2017)

TBC

TBC

- 1 Introduction
- 2 Feed-forward neural networks
- 3 Inference algorithms
- 4 Priors in weight space
- 5 Priors in function space**
- 6 Performance metrics

- 1 Introduction
- 2 Feed-forward neural networks
- 3 Inference algorithms
- 4 Priors in weight space
- 5 Priors in function space
- 6 Performance metrics**

- ▶ Predictive performance: ability of the model to give correct answers. Based on metrics (eg mean square error)
- ▶ Model calibration: assessing that the network is neither overconfident nor underconfident about its prediction. Requires using a test set.



- [1] Arsenii Ashukha, Alexander Lyzhov, Dmitry Molchanov, and Dmitry Vetrov. “Pitfalls of in-domain uncertainty estimation and ensembling in deep learning”. In: *International Conference on Learning Representations* (2020).
- [2] Andrew R Barron. “Approximation and estimation bounds for artificial neural networks”. In: *Machine Learning* 14.1 (1994), pp. 115–133.
- [3] Mikhail Belkin, Daniel Hsu, Siyuan Ma, and Soumik Mandal. “Reconciling modern machine-learning practice and the classical bias–variance trade-off”. In: *National Academy of Sciences* 116.32 (2019), pp. 15849–15854.
- [4] Vaggos Chatziafratis, Sai Ganesh Nagarajan, and Ioannis Panageas. “Better depth-width trade-offs for neural networks through the lens of dynamical systems”. In: *International Conference on Machine Learning*. 2020.

- [5] Vaggos Chatziafratis, Sai Ganesh Nagarajan, Ioannis Panageas, and Xiao Wang. “Depth-width trade-offs for ReLU networks via Sharkovsky’s theorem”. In: *International Conference on Learning Representations* (2020).
- [6] George Cybenko. “Approximation by superpositions of a sigmoidal function”. In: *Mathematics of Control, Signals and Systems* 2.4 (1989), pp. 303–314.
- [7] Ken-Ichi Funahashi. “On the approximate realization of continuous mappings by neural networks”. In: *Neural Networks* 2.3 (1989), pp. 183–192.
- [8] Xavier Glorot and Yoshua Bengio. “Understanding the difficulty of training deep feedforward neural networks”. In: *International Conference on Artificial Intelligence and Statistics*. 2010.
- [9] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. “Deep residual learning for image recognition”. In: *Computer Vision and Pattern Recognition*. 2016.

- [10] Matthias Hein, Maksym Andriushchenko, and Julian Bitterwolf. “Why ReLU networks yield high-confidence predictions far away from the training data and how to mitigate the problem”. In: *Computer Vision and Pattern Recognition*. 2019.
- [11] Sepp Hochreiter and Jürgen Schmidhuber. “Long short-term memory”. In: *Neural Computation* 9.8 (1997), pp. 1735–1780.
- [12] Kurt Hornik, Maxwell Stinchcombe, and Halbert White. “Multilayer feedforward networks are universal approximators”. In: *Neural Networks* 2.5 (1989), pp. 359–366.
- [13] Laurent Valentin Jospin, Wray Buntine, Farid Boussaid, Hamid Laga, and Mohammed Bennamoun. “Hands-on Bayesian Neural Networks—a Tutorial for Deep Learning Users”. In: *arXiv preprint arXiv:2007.06823* (2020).
- [14] Kimin Lee, Honglak Lee, Kibok Lee, and Jinwoo Shin. “Training confidence-calibrated classifiers for detecting out-of-distribution samples”. In: *International Conference on Learning Representations* (2018).

## References IV

- [15] Matthias Minderer, Josip Djolonga, Rob Romijnders, Frances Hubis, Xiaohua Zhai, Neil Houlsby, Dustin Tran, and Mario Lucic. “Revisiting the Calibration of Modern Neural Networks”. In: *Advances in Neural Information Processing Systems* (2021).
- [16] John Mitros and Brian Mac Namee. “On the validity of Bayesian neural networks for uncertainty estimation”. In: *arXiv preprint arXiv:1912.01530* (2019).
- [17] Yaniv Ovadia, Emily Fertig, Jie Ren, Zachary Nado, David Sculley, Sebastian Nowozin, Joshua V Dillon, Balaji Lakshminarayanan, and Jasper Snoek. “Can you trust your model’s uncertainty? Evaluating predictive uncertainty under dataset shift”. In: *Advances in Neural Information Processing Systems* (2019).
- [18] Ben Poole, Subhaneil Lahiri, Maithra Raghu, Jascha Sohl-Dickstein, and Surya Ganguli. “Exponential expressivity in deep neural networks through transient chaos”. In: *International Conference on Neural Information Processing Systems*. 2016.

- [19] Herbert Robbins and Sutton Monro. “A stochastic approximation method”. In: *The Annals of Mathematical Statistics* (1951), pp. 400–407.
- [20] David E Rumelhart, Geoffrey E Hinton, and Ronald J Williams. “Learning representations by back-propagating errors”. In: *Nature* 323.6088 (1986), pp. 533–536.
- [21] Samuel S Schoenholz, Justin Gilmer, Surya Ganguli, and Jascha Sohl-Dickstein. “Deep information propagation”. In: *International Conference on Learning Representations*. 2017.
- [22] Shiliang Sun, Zehui Cao, Han Zhu, and Jing Zhao. “A survey of optimization methods from a machine learning perspective”. In: *IEEE Transactions on Cybernetics* 50.8 (2019), pp. 3668–3681.
- [23] Andrew Gordon Wilson, Zhiting Hu, Ruslan Salakhutdinov, and Eric P Xing. “Deep kernel learning”. In: *International Conference on Artificial Intelligence and Statistics*. 2016.