

BML lecture #7: Generative priors

<http://github.com/rbardenet/bml-course>

Gabriel V Cardoso

`gabriel.victorino_cardoso@minesparis.psl.eu`

Ecole des Mines, Univ. PSL, France



Generative priors: Using generative models to solve (Bayesian) inverse problems.

- ▶ Part I: A fast introduction to score based generative model (SGM)
- ▶ Part II: How to use SGM as prior for inverse problems

The original miracle!



Figure: $X_0 \sim p_D$.



Figure: $X_\sigma = X_0 + \sigma Z$.¹

- ▶ D_{θ_0} untrained CNN,
- ▶ $L_\sigma(\theta) = \|D_\theta(X_\sigma) - X_\sigma\|^2$, $L_0(\theta) = \|D_\theta(X_\sigma) - X_0\|^2$,
- ▶ $\theta_t = \theta_{t-1} - \nabla L_\sigma(\theta_{t-1})$

What happens to $\{L_\sigma(\theta_t)\}_{t=1}^N$ and $\{L_0(\theta_t)\}_{t=1}^N$?

¹ $Z \sim \mathcal{N}(0, I)$, $X_0 \perp Z$.

The original miracle!



Figure: $X_0 \sim p_{\mathcal{D}}$.



Figure: $X_{\sigma} = X_0 + \sigma Z$.¹

- ▶ D_{θ_0} untrained CNN,
- ▶ $L_{\sigma}(\theta) = \|D_{\theta}(X_{\sigma}) - X_{\sigma}\|^2$, $L_0(\theta) = \|D_{\theta}(X_{\sigma}) - X_0\|^2$,
- ▶ $\theta_t = \theta_{t-1} - \nabla L_{\sigma}(\theta_{t-1})$

What happens to $\{L_{\sigma}(\theta_t)\}_{t=1}^N$ and $\{L_0(\theta_t)\}_{t=1}^N$?

¹ $Z \sim \mathcal{N}(0, I)$, $X_0 \perp Z$.

The original miracle!



Figure: $X_0 \sim p_D$.






Figure: $X_\sigma = X_0 + \sigma Z$.¹

- ▶ D_{θ_0} untrained CNN,
- ▶ $L_\sigma(\theta) = \|D_\theta(X_\sigma) - X_\sigma\|^2$, $L_0(\theta) = \|D_\theta(X_\sigma) - X_0\|^2$,
- ▶ $\theta_t = \theta_{t-1} - \nabla L_\sigma(\theta_{t-1})$

What happens to $\{L_\sigma(\theta_t)\}_{t=1}^N$ and $\{L_0(\theta_t)\}_{t=1}^N$?

¹ $Z \sim \mathcal{N}(0, I)$, $X_0 \perp Z$.

The original miracle ².



²This experiment is inspired by Ulyanov, Vedaldi, and Lempitsky, 2018

How to can we take advantage of extremely good denoisers to create a generative model?

Lemma (Vincent, 2011)

$$\mathbb{E}[X_0 | X_\sigma = x] = x + \sigma^2 \nabla \log p_\sigma(x), \quad (1)$$

where $p_\sigma = \text{Law}(X_\sigma)$.

$$0 < \sigma_1 < \dots < \sigma_T.$$

$$\theta \in \arg \min \sum_{t=1}^N \gamma_t^2 \mathbb{E} [\|D_\theta(X_0 + \sigma_t Z, \sigma_t) - X_0\|^2] \quad (2)$$

Sequential Langevin Algorithm:

$X_T \sim \mathcal{N}(0, \sigma_T^2 \mathbf{I})$, define $s_\theta(x, \sigma) = \sigma^{-2}(D_\theta(x, \sigma) - x)$.

For $t = T - 1, \dots, 1$

- ▶ $X_{0,t} = X_{K,t+1}$
- ▶ $X_{k,t} = X_{k-1,t} + \eta s_\theta(X_{k-1,t}, \sigma) + \sqrt{2\eta} Z_{k,t}.$

$$0 < \sigma_1 < \cdots < \sigma_T.$$

$$\theta \in \arg \min \sum_{t=1}^N \gamma_t^2 \mathbb{E} [\|D_\theta(X_0 + \sigma_t Z, \sigma_t) - X_0\|^2] \quad (2)$$

Sequential Langevin Algorithm:

$X_T \sim \mathcal{N}(0, \sigma_T^2 \mathbf{I})$, define $s_\theta(x, \sigma) = \sigma^{-2}(D_\theta(x, \sigma) - x)$.

For $t = T - 1, \dots, 1$

- ▶ $X_{0,t} = X_{K,t+1}$
- ▶ $X_{k,t} = X_{k-1,t} + \eta s_\theta(X_{k-1,t}, \sigma) + \sqrt{2\eta} Z_{k,t}.$

$$0 < \sigma_1 < \dots < \sigma_T.$$

$$\theta \in \arg \min \sum_{t=1}^N \gamma_t^2 \mathbb{E} [\|D_\theta(X_0 + \sigma_t Z, \sigma_t) - X_0\|^2] \quad (2)$$

Sequential Langevin Algorithm:

$X_T \sim \mathcal{N}(0, \sigma_T^2 \mathbf{I})$, define $s_\theta(x, \sigma) = \sigma^{-2}(D_\theta(x, \sigma) - x)$.

For $t = T - 1, \dots, 1$

- ▶ $X_{0,t} = X_{K,t+1}$
- ▶ $X_{k,t} = X_{k-1,t} + \eta s_\theta(X_{k-1,t}, \sigma) + \sqrt{2\eta} Z_{k,t}.$

Consider the following Markov chain:

- 1 $X_0 \sim p_{\mathcal{D}},$
- 2 $X_t \sim X_{t-1} + \sqrt{\sigma_t^2 - \sigma_{t-1}^2}.$

The joint law of this chain is

$$p_{0:T}(x_{0:T}) = p_{\mathcal{D}}(x_0) \prod_{t=1}^T \overbrace{\mathcal{N}(x_t; x_{t-1}, (\sigma_t^2 - \sigma_{t-1}^2)\mathbf{I})}^{:= p_{t|t-1}(x_t|x_{t-1})}. \quad (3)$$

It holds that:

- 1 $p_{t|0}(x_t|x_0) = \mathcal{N}(x_0, \sigma_t^2\mathbf{I}),$
- 2 $p_{t|0,t+1}(x_t|x_0, x_{t+1}) = \mathcal{N}(x_t, x_0 + \frac{\sigma_t}{\sigma_{t+1}}(x_{t+1} - x_0), \frac{\sigma_t^2}{\sigma_{t+1}^2}(\sigma_{t+1}^2 - \sigma_t^2)\mathbf{I})$

Consider the following Markov chain:

- 1 $X_0 \sim p_{\mathcal{D}},$

- 2 $X_t \sim X_{t-1} + \sqrt{\sigma_t^2 - \sigma_{t-1}^2}.$

The joint law of this chain is

$$p_{0:T}(x_{0:T}) = p_{\mathcal{D}}(x_0) \prod_{t=1}^T \overbrace{\mathcal{N}(x_t; x_{t-1}, (\sigma_t^2 - \sigma_{t-1}^2)\mathbf{I})}^{:= p_{t|t-1}(x_t|x_{t-1})}. \quad (3)$$

It holds that:

- 1 $p_{t|0}(x_t|x_0) = \mathcal{N}(x_0, \sigma_t^2\mathbf{I}),$

- 2 $p_{t|0,t+1}(x_t|x_0, x_{t+1}) = \mathcal{N}(x_t, x_0 + \frac{\sigma_t}{\sigma_{t+1}^2}(x_{t+1} - x_0), \frac{\sigma_t^2}{\sigma_{t+1}^2}(\sigma_{t+1}^2 - \sigma_t^2)\mathbf{I})$

Consider the following Markov chain:

- 1 $X_0 \sim p_{\mathcal{D}},$

- 2 $X_t \sim X_{t-1} + \sqrt{\sigma_t^2 - \sigma_{t-1}^2}.$

The joint law of this chain is

$$p_{0:T}(x_{0:T}) = p_{\mathcal{D}}(x_0) \prod_{t=1}^T \overbrace{\mathcal{N}(x_t; x_{t-1}, (\sigma_t^2 - \sigma_{t-1}^2)\mathbf{I})}^{:= p_{t|t-1}(x_t|x_{t-1})}. \quad (3)$$

It holds that:

- 1 $p_{t|0}(x_t|x_0) = \mathcal{N}(x_0, \sigma_t^2\mathbf{I}),$

- 2 $p_{t|0,t+1}(x_t|x_0, x_{t+1}) = \mathcal{N}(x_t, x_0 + \frac{\sigma_t}{\sigma_{t+1}^2}(x_{t+1} - x_0), \frac{\sigma_t^2}{\sigma_{t+1}^2}(\sigma_{t+1}^2 - \sigma_t^2)\mathbf{I})$

We want to reverse the noising process. Consider

$$q_{0:T}^{\theta}(x_{0:T}) := \mathcal{N}(x_T; \mu_T, \eta_T^2 \mathbf{I}) \prod_{t=0}^{T-1} \mathcal{N}(x_t; \mu_{t,\theta}(x_{t+1}), \eta_{t+1}^2 \mathbf{I}) . \quad (4)$$

Our goal is to

$$\arg \min_{\theta, \eta_{0:T}^2, \mu_T} D_{\text{KL}}(p_{0:T} || q_{0:T}^{\theta}) . \quad (5)$$

Before advancing, note that $p_{0:T}$ (as it is a Markov chain) admits the following decomposition

$$p_{0:T}(x_{0:T}) = p_T(x_T) \prod_{t=0}^{T-1} p_{t|t+1}(x_t | x_{t+1}) .$$

Unfortunately, each $p_{t|t+1}$ is intractable for general p_D .

We want to reverse the noising process. Consider

$$q_{0:T}^{\theta}(x_{0:T}) := \mathcal{N}(x_T; \mu_T, \eta_T^2 \mathbf{I}) \prod_{t=0}^{T-1} \mathcal{N}(x_t; \mu_{t,\theta}(x_{t+1}), \eta_{t+1}^2 \mathbf{I}) . \quad (4)$$

Our goal is to

$$\arg \min_{\theta, \eta_{0:T}^2, \mu_T} \mathrm{D}_{\mathrm{KL}}(p_{0:T} \| q_{0:T}^{\theta}) . \quad (5)$$

Before advancing, note that $p_{0:T}$ (as it is a Markov chain) admits the following decomposition

$$p_{0:T}(x_{0:T}) = p_T(x_T) \prod_{t=0}^{T-1} p_{t|t+1}(x_t | x_{t+1}) .$$

Unfortunately, each $p_{t|t+1}$ is intractable for general p_D .

We want to reverse the noising process. Consider

$$q_{0:T}^{\theta}(x_{0:T}) := \mathcal{N}(x_T; \mu_T, \eta_T^2 \mathbf{I}) \prod_{t=0}^{T-1} \mathcal{N}(x_t; \mu_{t,\theta}(x_{t+1}), \eta_{t+1}^2 \mathbf{I}) . \quad (4)$$

Our goal is to

$$\arg \min_{\theta, \eta_{0:T}^2, \mu_T} D_{\text{KL}}(p_{0:T} \| q_{0:T}^{\theta}) . \quad (5)$$

Before advancing, note that $p_{0:T}$ (as it is a Markov chain) admits the following decomposition

$$p_{0:T}(x_{0:T}) = p_T(x_T) \prod_{t=0}^{T-1} p_{t|t+1}(x_t | x_{t+1}) .$$

Unfortunately, each $p_{t|t+1}$ is intractable for general $p_{\mathcal{D}}$.

$$\begin{aligned}
 D_{\text{KL}}(p_{0:T} || q_{0:T}^\theta) &= \\
 \mathbb{E}_{p_{0:T}} \left[\log \left(\frac{p_T(x_T) \prod_{t=0}^{T-1} p_{t|t+1}(x_t | x_{t+1})}{\mathcal{N}(x_T; \mu_T, \eta_T^2 \mathbf{I}) \prod_{t=0}^{T-1} \mathcal{N}(x_t; \mu_{t,\theta}(x_{t+1}), \eta_{t+1}^2 \mathbf{I})} \right) \right] \\
 &= D_{\text{KL}}(p_T || \mathcal{N}(\mu_T, \eta_T^2 \mathbf{I})) \\
 &\quad + \sum_{t=1}^{T-1} \mathbb{E}_{p_{0:T}} [D_{\text{KL}}(p_{t|t+1}(\cdot | X_{t+1}) || \mathcal{N}(\mu_{t,\theta}(X_{t+1}), \eta_{t+1}^2 \mathbf{I}))] \\
 &\quad + \mathbb{E}_{p_{0:T}} [D_{\text{KL}}(p_{\mathcal{D}} || \mathcal{N}(\mu_{0,\theta}(X_1), \eta_{t+1}^2 \mathbf{I}))] .
 \end{aligned}$$

We now use the following property:

Lemma

If $\lambda \in \mathcal{P}_2$, then

$$(\mu_\lambda, \Sigma_\lambda) \in \arg \min D_{\text{KL}}(\lambda || \mathcal{N}(\mu, \Sigma)) ,$$

where $\mu_\lambda = \mathbb{E}_\lambda[X]$ and $\Sigma = \mathbb{V}_\lambda[X] = \mathbb{E}_\lambda[(X - \mu_\lambda)(X - \mu_\lambda)^T]$.

$$\begin{aligned}
 D_{\text{KL}}(p_{0:T} || q_{0:T}^\theta) &= \\
 \mathbb{E}_{p_{0:T}} \left[\log \left(\frac{p_T(x_T) \prod_{t=0}^{T-1} p_{t|t+1}(x_t | x_{t+1})}{\mathcal{N}(x_T; \mu_T, \eta_T^2 \mathbf{I}) \prod_{t=0}^{T-1} \mathcal{N}(x_t; \mu_{t,\theta}(x_{t+1}), \eta_{t+1}^2 \mathbf{I})} \right) \right] \\
 &= D_{\text{KL}}(p_T || \mathcal{N}(\mu_T, \eta_T^2 \mathbf{I})) \\
 &\quad + \sum_{t=1}^{T-1} \mathbb{E}_{p_{0:T}} [D_{\text{KL}}(p_{t|t+1}(\cdot | X_{t+1}) || \mathcal{N}(\mu_{t,\theta}(X_{t+1}), \eta_{t+1}^2 \mathbf{I}))] \\
 &\quad + \mathbb{E}_{p_{0:T}} [D_{\text{KL}}(p_{\mathcal{D}} || \mathcal{N}(\mu_{0,\theta}(X_1), \eta_{t+1}^2 \mathbf{I}))] .
 \end{aligned}$$

We now use the following property:

Lemma

If $\lambda \in \mathcal{P}_2$, then

$$(\mu_\lambda, \Sigma_\lambda) \in \arg \min D_{\text{KL}}(\lambda || \mathcal{N}(\mu, \Sigma)) ,$$

where $\mu_\lambda = \mathbb{E}_\lambda[X]$ and $\Sigma = \mathbb{V}_\lambda[X] = \mathbb{E}_\lambda[(X - \mu_\lambda)(X - \mu_\lambda)^T]$.

$$D_{\text{KL}}(p_T || \mathcal{N}(\mu_T, \eta_T^2 \mathbf{I}))$$

$$\mathbb{E}_{\rho_{0:T}} \left[D_{\text{KL}} \left(p_{t|t+1}(\cdot | X_{t+1}) || \mathcal{N}(\mu_{t,\theta}(X_{t+1}), \eta_{t+1}^2 \mathbf{I}) \right) \right]$$

$$\mathbb{E}_{p_{0:T}} \left[D_{\text{KL}} \left(p_{\mathcal{D}} \middle| \middle| \mathcal{N}(\mu_{0,\theta}(X_1), \eta_{t+1}^2) \right) \right]$$

Algorithm:

1 Sample $X_T \sim \mathcal{N}(0, \sigma_T^2)$,

2 Sample for $t = T - 1, \dots, 1$,

$$X_t \sim \mathcal{N} \left(D_\theta(X_{t+1}, \sigma_{t+1}) + \frac{\sigma_t^2}{\sigma_{t+1}^2} (X_{t+1} - D_\theta(X_{t+1}, \sigma_{t+1})), \frac{\sigma_t^2}{\sigma_{t+1}^2} (\sigma_{t+1}^2 - \sigma_t^2) \mathbf{I} \right) .$$

Often in the literature, you see the following noising process for $\beta_t \in (0, 1)$:

$$Y_t = \sqrt{1 - \beta_t} Y_{t-1} + \sqrt{\beta_t} Z_t . \quad (6)$$

It has the nice property that $\mathbb{E}[Y_t] = 0$ and $\mathbb{V}[Y_t] = 1$ (Thus variance preserving) and that

$$p(y_t|y_0) = \mathcal{N}(\sqrt{\alpha_t} y_0, (1 - \alpha_t) \mathbf{I})$$

where $\alpha_t = \prod_{k \leq t} (1 - \beta_k)$.

It can be obtained from our process by defining $Y_t = X_t / \sqrt{\sigma_t^2 + 1}$.

Furthermore, knowing the score of X_t can give access to the score of Y_t ! (The two are equivalent).

Often in the literature, you see the following noising process for $\beta_t \in (0, 1)$:

$$Y_t = \sqrt{1 - \beta_t} Y_{t-1} + \sqrt{\beta_t} Z_t . \quad (6)$$

It has the nice property that $\mathbb{E}[Y_t] = 0$ and $\mathbb{V}[Y_t] = 1$ (Thus variance preserving) and that

$$p(y_t|y_0) = \mathcal{N}(\sqrt{\alpha_t} y_0, (1 - \alpha_t) \mathbf{I})$$

where $\alpha_t = \prod_{k \leq t} (1 - \beta_k)$.

It can be obtained from our process by defining $Y_t = X_t / \sqrt{\sigma_t^2 + 1}$.

Furthermore, knowing the score of X_t can give access to the score of Y_t ! (The two are equivalent).

Algorithm:

- 1 Sample $X_T \sim \mathcal{N}(0, \mathbf{I})$,
- 2 Sample for $t = T - 1, \dots, 1$,

$$X_t \sim \mathcal{N} \left(\frac{1}{\sqrt{1 - \beta_{t+1}}} (X_{t+1} - \frac{\beta_{t+1}}{\sqrt{1 - \alpha_t}} \varepsilon_{\theta}(X_{t+1}, \sqrt{\alpha_{t+1}})), \frac{\sigma_t^2}{\sigma_{t+1}^2} (\sigma_{t+1}^2 - \sigma_t^2) \mathbf{I} \right),$$

where the **epsilon net** ε_{θ} can be linked to our denoising network by:

$$D_{\theta}(x, \sigma(\alpha_t)) = \frac{1}{\sqrt{\beta_t}} \left(x - \frac{\beta_t}{\sqrt{1 - \alpha_t}} \varepsilon_{\theta}(X_{t+1}, \sqrt{\alpha_t}) \right),$$

where $\sigma(\alpha_t) = \sqrt{\frac{1 - \alpha_t}{\alpha_t}}$.

Remember our inverse problem posterior:

$$p_0(x|y) = \frac{\ell(y|x)p_0(x)}{\int \ell(y|\tilde{x})p_0(\tilde{x})d\tilde{x}} . \quad (7)$$

Can we use a diffusion model to sample from $p_0(x|y)$?

Remember our inverse problem posterior:

$$p_0(x|y) = \frac{\ell(y|x)p_0(x)}{\int \ell(y|\tilde{x})p_0(\tilde{x})d\tilde{x}} . \quad (7)$$

Can we use a diffusion model to sample from $p_0(x|y)$?

We start by defining

$$p_t(x|y) = \int p_0(x_0|y) \mathcal{N}(x; x_0, \sigma_t^2 \mathbf{I}) dx_0 . \quad (8)$$

What is the score?

Chung et al., 2023 proposes the following approximation

$$\nabla \log p_t(x|y) = \nabla \log p_t(x) + \zeta_t \nabla \log l(y| \underbrace{\mathbb{E}[X_0|X_t = x_t]}_{x + \sigma^2 \nabla \log p_t(x)}) \quad (9)$$

In practice, for linear problems (i.e.) $\log l(y|x) = -\frac{1}{2\sigma_y^2} \|y - Ax\|^2 + c$, the authors suggests setting $\zeta_t = \tilde{\zeta} \frac{2\sigma_y^2}{\|y - Ax\|}$ where $\tilde{\zeta}$ is a tunable constant (typically 1 or 0.1).

Chung et al., 2023 proposes the following approximation

$$\nabla \log p_t(x|y) = \nabla \log p_t(x) + \zeta_t \nabla \log l(y| \underbrace{\mathbb{E}[X_0|X_t = x_t]}_{x + \sigma^2 \nabla \log p_t(x)}) \quad (9)$$

In practice, for linear problems (i.e.) $\log l(y|x) = -\frac{1}{2\sigma_y^2} \|y - Ax\|^2 + c$, the authors suggests setting $\zeta_t = \tilde{\zeta} \frac{2\sigma_y^2}{\|y - Ax\|}$ where $\tilde{\zeta}$ is a tunable constant (typically 1 or 0.1).

- [1] Hyungjin Chung et al. “Diffusion Posterior Sampling for General Noisy Inverse Problems”. In: *The Eleventh International Conference on Learning Representations*. 2023.
- [2] Jonathan Ho, Ajay Jain, and Pieter Abbeel. “Denoising Diffusion Probabilistic Models”. In: *Advances in Neural Information Processing Systems*. Vol. 33. 2020, pp. 6840–6851.
- [3] Yang Song and Stefano Ermon. “Generative Modeling by Estimating Gradients of the Data Distribution”. In: *Advances in neural information processing systems* 32 (2019).
- [4] Dmitry Ulyanov, Andrea Vedaldi, and Victor Lempitsky. “Deep Image Prior”. In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. June 2018.
- [5] Pascal Vincent. “A Connection Between Score Matching and Denoising Autoencoders”. In: *Neural Computation* 23.7 (2011), pp. 1661–1674. DOI: 10.1162/NECO_a_00142.